📝 Notes

Apr 10, 2026

UX Design for Marc - Command Center

Invited Supreo Mukhopadhyay Neha Bharati Tanmay Jyot Nasahn Sheppard Muddu Sudhakar Aditi Sharma Alex Kahng Shantanu Singh

Attachments UX Design for Marc - Command Center Notes by Gemini

Meeting records Transcript

Summary

Visualization improvements and Agent Force implementation strategies defined for agent management via a decision theater tool.

Visualization and Readability Fixes
Visual accuracy issues in maps required immediate correction. Teams finalized increasing font sizes to improve overall interface readability.

Agent Force Panel Integration
Developers will implement the Agent Force panel using a conversational interface for managing agent policies. This feature replaces complex correction panels with streamlined human interventions.

Strategic Product Focus
The product evolved into a decision theater tool prioritizing security and production data management. Developers decided to prioritize traffic control and human in the loop features.

Next steps

- [ ] [Aditi Sharma] Update Map: Redo map visualization; remove dots from ocean area. Ensure accurate detail at current zoom level.
- [ ] [Alex Kahng] Design Policy: Develop a concept for policy structure, potentially using an MD file format. Ensure agents adhere to established policy guidelines.
- [ ] [Aditi Sharma] Integrate Policy: Add policy document view capability to the agent screen. Ensure ability to click and display the associated document.
- [ ] [Aditi Sharma] Increase Font: Increase all visualization text and character font sizes. Adjust display to improve readability for projection.
- [ ] [Aditi Sharma] Research Agent: Conduct research into appropriate fixes applied to agents based on policy breaches.
- [ ] [Aditi Sharma] Create Agentforce: Build workflow story for side panel functionality. Ensure behavior mirrors Claude Code conversational experience.
- [ ] [Alex Kahng] Review Story: Review Aditis Agent Force story documentation. Provide thumbs up or thumbs down feedback within the Google doc.
- [ ] [Aditi Sharma] Record Demo: Create a video recording of the interactive site demo. Use a large monitor during recording to ensure improved screen resolution.
- [ ] [Muddu Sudhakar] Discuss Presentation: Meet with Mayank to discuss and work on the final presentation deck.

Details

Did the screenshots in this section make your notes better or worse?

- **Review of Current Visualization and Map Accuracy**: The discussion began with feedback on the current visualization, with Muddu Sudhakar noting that the map was confusing and "messed up," particularly mentioning the number of green dots in the ocean. Alex Kahng agreed that this part was "odd". Muddu Sudhakar advised that if the visualization is intended to convey detail, it must be exact, and noted the need to fix the visual issues before presenting it to Mark.

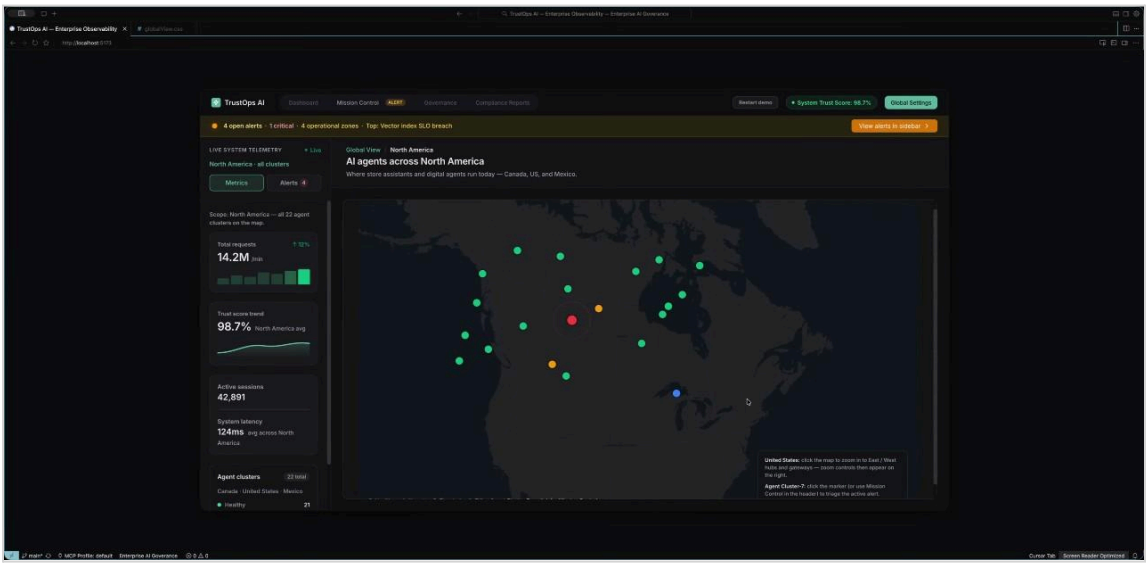

- **Agent Monitoring and Alert System**: Aditi Sharma explained that the screen shows data and alerts on the sidebar, allowing users to zoom in to see specific hubs (like US East and US West) and the deployed agents. They indicated that the screen should show where in the workflow an agent is breaking and what the live traffic is, referencing Alex Kahng's feedback about visually showing the blast radius. The ability for a user to redeploy an agent after solving a problem from the same screen was also highlighted (00:00:36).

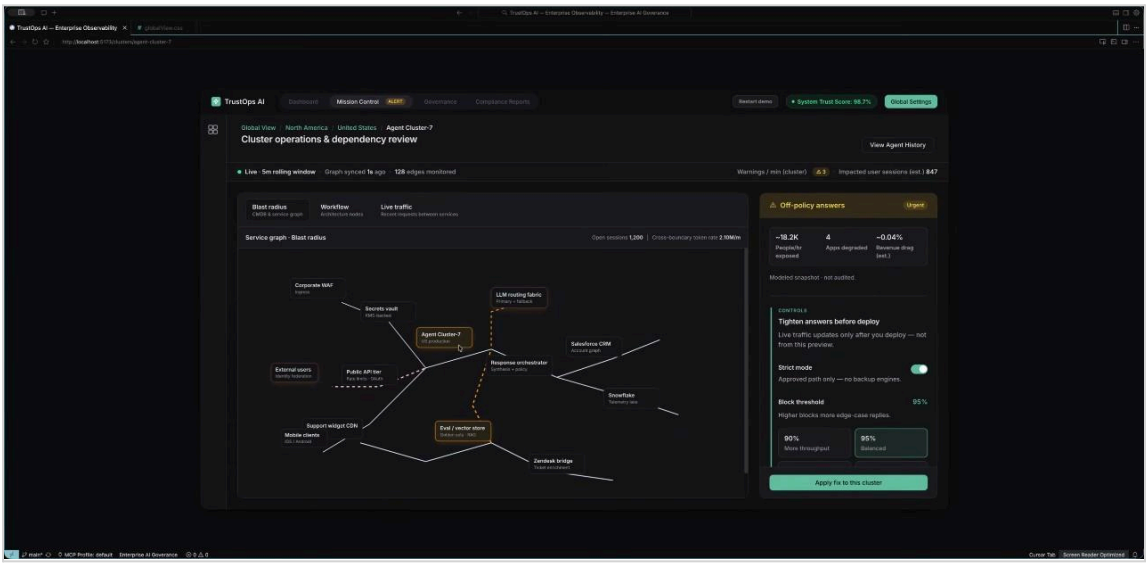

- **"Human in the Loop" Definition and Policy Visualization**: Muddu Sudhakar requested clarification on "human in the loop," which Aditi Sharma clarified is an option to intervene if an agent is behaving poorly. Muddu Sudhakar emphasized

the critical need to show policy adherence, asking how to determine if an agent is following its defined policy and when it violates policy guidelines (00:01:42). Muddu Sudhakar suggested that the policy should be written in a structured format, like an MD file, and be viewable by clicking on the agent (00:02:31).

- **Font Size and Readability Issues**: Muddu Sudhakar repeatedly stressed that the font size across the entire view was too small, making characters difficult to read, especially when projected, and suggested increasing the font size by "a couple of order of magnitude" (00:03:20). Alex Kahng confirmed that the text was too small and that they could not read it. Aditi Sharma accepted the feedback and committed to changing the font size (00:04:08).
- **Remediation Steps and Agent Force Implementation**: Alex Kahng suggested replacing the current detailed correction panel with a feature called "Agent Force" that uses a conversational flow for policy changes (00:04:08). Alex Kahng emphasized that remediation should involve showing the policy clause that was breached via a tooltip and simplifying the remediation to changing the evaluation set or adding cases, followed by a retraining process (00:05:08). They concluded that the human intervention should be high-level, focusing on three actions: "add to data set," "retrain agent," and "launch agent," keeping the internal configurations as a black box (00:06:18) (00:08:47).

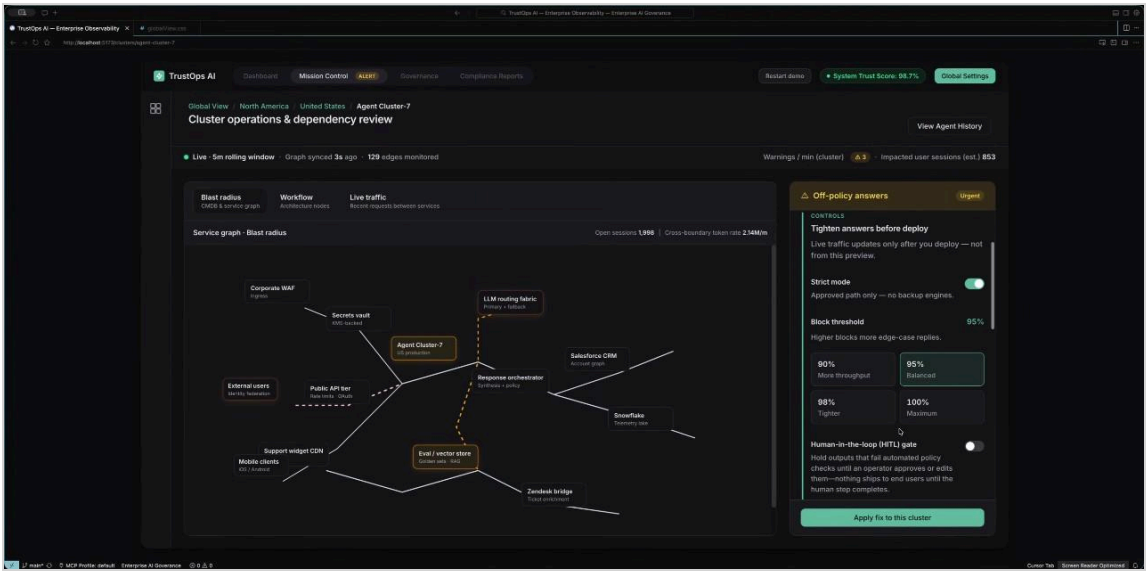

- **Implementation of Conversational Agent Force Panel**: Alex Kahng proposed that the right-side panel should be modeled after cloud code, branded as Agent Force, offering a conversational interface for making changes like updating evals and configuration in a black-box manner (00:07:33). Aditi Sharma confirmed they

would create a workflow where the side panel behaves like cloud code, allowing the user to tell the agent to add data sets, run evals, and then ask if the agent should be launched again, ensuring it is "eval driven" (00:09:48).

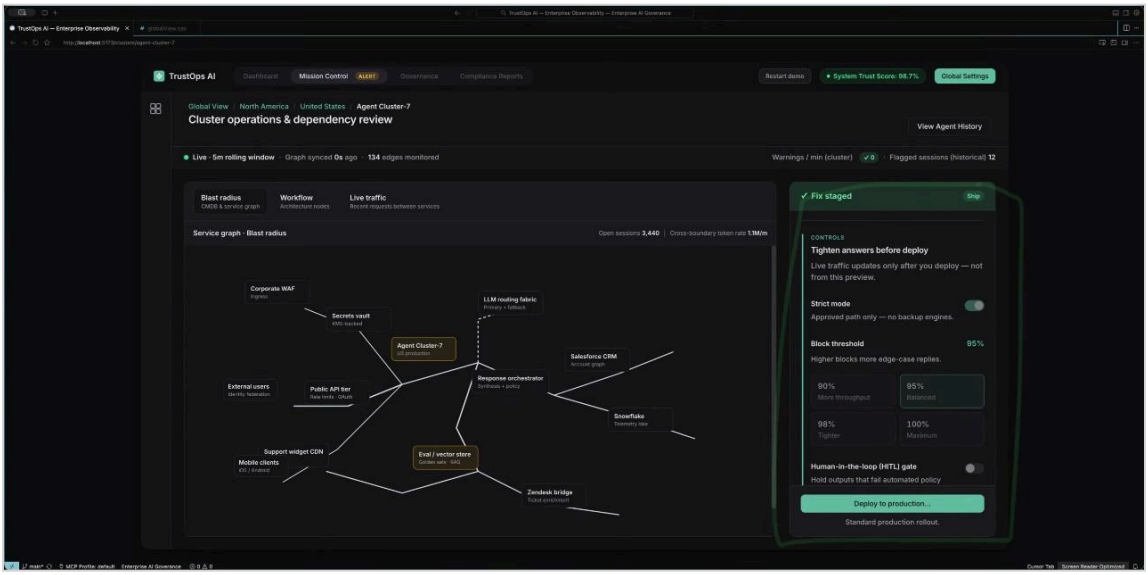

- **Action-Oriented Interaction Patterns**: The team discussed the need for different interaction patterns beyond just agent life cycle management, including high-stakes emergency situations (00:10:33). Alex Kahng suggested interactions like the ability to stop an agent ("kill switch") or swap agents, noting that their approach should be more action-oriented than competing products, which are primarily focused on observability (00:11:40). Muddu Sudhakar highlighted that competitors are focused on observability and evaluations, while their team needs to focus on red teaming, security, and using production data for evolution (00:12:37).
- **Controlling Agent Traffic and the Role of the User**: The discussion shifted from a "kill switch" to the ability to seamlessly reroute traffic away from a failing agent (00:18:07). Alex Kahng suggested routing traffic visually or conversationally using the right-side panel, for example, by asking Agent Force to "stop routing any traffic to agent cluster 7 or reduce traffic to 2%" (00:18:55). The group acknowledged that the user persona is still evolving, but Alex Kahng suggested thinking of the future user as a young "savant" who manages these agents as an abstract view, which Muddu Sudhakar agreed with (00:20:52).

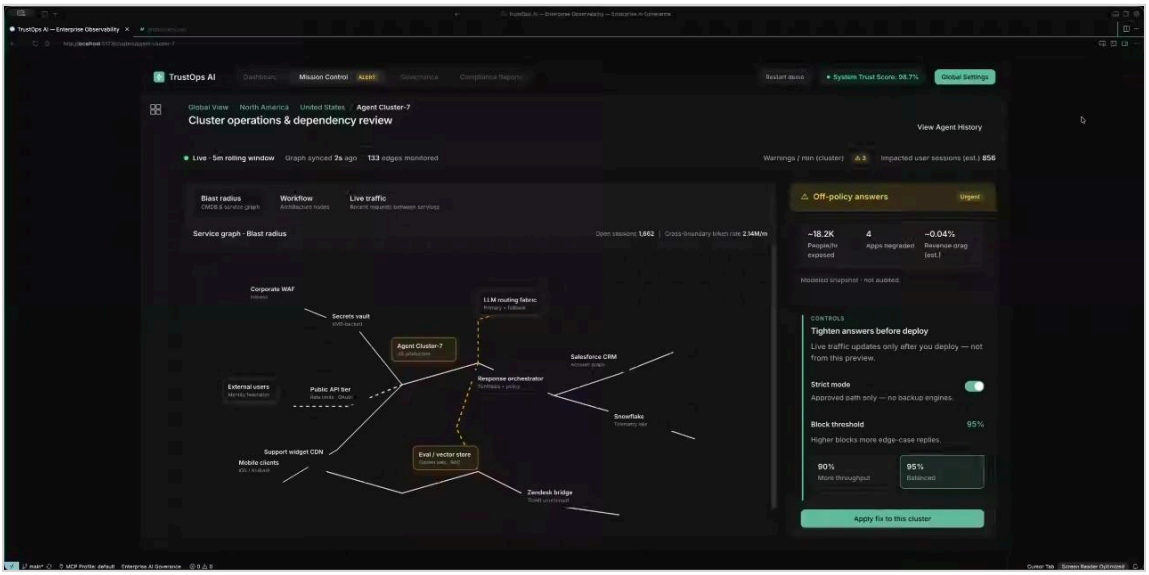

- **Product Focus and Evolution of Agent Management**: The purpose of the product was clarified as being a "decision theater" tool for making practical decisions to manage agents in production (00:22:25). Muddu Sudhakar explained the evolution of customer needs: first, getting value and utility from agents, then worrying about uptime and performance (observability/control tower), and finally, mitigating risk from rogue agents (security) (00:23:08). Aditi Sharma summarized the next steps: making the map more accurate, increasing the font size, and making the side panel more conversational and agent force-driven to cover stories like "human in the loop" (00:25:29) (00:31:59).
- **Final Deliverables and Meeting Preparation**: Muddu Sudhakar requested that the final deliverable be a video for a meeting with Mark, and Alex Kahng confirmed the need for both the interactive site and the video (00:28:53). Muddu Sudhakar also asked Aditi Sharma to record the video without an 11 Labs voiceover. Shantanu Singh confirmed that the interactive version of the product is hosted on Heroku (00:29:30). Aditi Sharma finalized the changes, including a 4x font increase and a conversational Agent Force panel showing the human-in-the-loop and traffic reduction stories (00:27:03) (00:31:59).

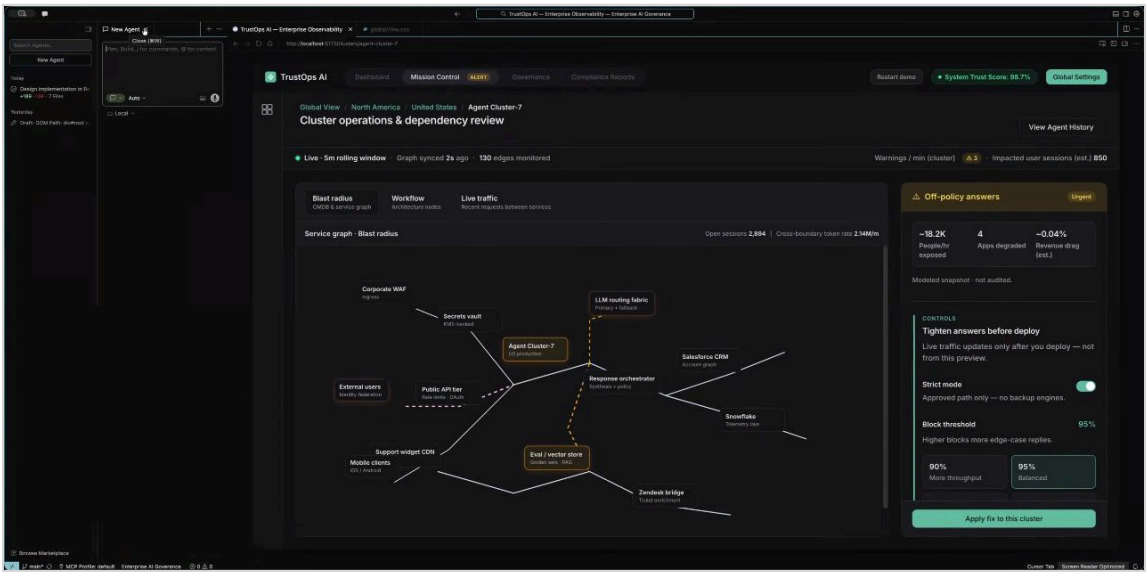

You should review Gemini's notes to make sure they're accurate. [Get tips and learn how Gemini takes notes]

*How is the quality of **these specific notes?** [Take a short survey] to let us know your feedback, including how helpful the notes were for your needs.*

📖 Transcript

Apr 9, 2026

UX Design for Marc - Command Center - Transcript

00:00:00

Muddu Sudhakar: Man, you have different visualization. You can tell everything in life.
Aditi Sharma: Okay, sure. No,
Alex Kahng: Yeah.
Muddu Sudhakar: No.
Aditi Sharma: I'll take that feedback.
Muddu Sudhakar: Fix it. This won't go well. This like people don't even know what this is,
Aditi Sharma: Yeah.
Muddu Sudhakar: where we are, nothing. So, and why are so many green dots in the ocean? It's like messed up.
Aditi Sharma: Okay. Okay,
Muddu Sudhakar: Don't put any dots in the ocean,
Aditi Sharma: then I Right.
Alex Kahng: Is it that
Muddu Sudhakar: first of all.
Aditi Sharma: Right. Right right right. Um,
Alex Kahng: part?
Aditi Sharma: yeah. Uh
Muddu Sudhakar: What is it, Alex?
Alex Kahng: Yeah, that part's odd.
Muddu Sudhakar: Yeah,
Aditi Sharma: so
Muddu Sudhakar: this like this looks like really messed up shot. If you want to convey this, make it exact in detail where possible at this level. Okay, next one.
Aditi Sharma: okay.
Muddu Sudhakar: Cuz I'm going to play this thing in front of Mark.

00:00:36

Aditi Sharma: So
Muddu Sudhakar: So just don't do like I don't want to lose him at this screen itself.
Aditi Sharma: right right. So um all the data you can also see alerts on the sidebar here or you can just like zoom in and you can see like uh US east hub and US west hub and the global gateway and all these agents that are deployed in those zones and then this is the particular agent which has some sort of like there is a disconnect here either it's uh failing at routing and then you can just look at where in the exact um workflow it's breaking what is the live traffic here and this sort of is the you know the feedback that we got from Alex was the blast radius needs to be um shown visually and one of the feedback that we got from Alex yesterday was that the user should be able to redeploy the agent after solving the problem from this screen itself. So you can like whatever the problem is like you can decide what threshold you want and you want human in the loop review and then you just like apply it and

00:01:42

Muddu Sudhakar: What is mean human in the loop? You're going too fast.
Aditi Sharma: uh uh yeah for the agent like
Muddu Sudhakar: What do you mean by human in the loop? It's for the agent.
Aditi Sharma: for example if an agent is behaving is behaving really badly I have an option of uh going uh
Muddu Sudhakar: Okay. You force a human into the loop.
Aditi Sharma: taking it.
Muddu Sudhakar: Yeah.
Aditi Sharma: Yeah. Yeah. So from here I can just like deploy it um this agent after those corrections from here
Muddu Sudhakar: Right. I think one thing it's missing here. Alex feel free to jump in to say see what I need to know is green is okay.
Aditi Sharma: itself.
Muddu Sudhakar: Is green okay from how their activity or is it also okay from the policy going forward what going to happen every agent will have a policy. What you're not showing me is which agents are actually following the policy they define they uh they have associated with. I should be able to click next to an agent and show its policy to see has it met all its policy guidelines.

00:02:31

Muddu Sudhakar: When is did it violate its policy guidelines? You know what I'm saying? The red,
Aditi Sharma: So in
Muddu Sudhakar: you see what I'm saying? Policy should be written as a do some kind of a structure like an MD file itself.
Aditi Sharma: Yeah.
Muddu Sudhakar: Alex, think of some kind of an MD file for policy and agent should follow that policy. You know what I'm saying?
Alex Kahng: Yep.
Muddu Sudhakar: Okay.
Alex Kahng: Yep.
Aditi Sharma: So right now in this story we're just going after that agent is not giving policy uh adjacent answers. They're going like they're giving off policy answers.
Muddu Sudhakar: Right.
Aditi Sharma: So what you mean is like there should
Muddu Sudhakar: But where is the policy though? I want to see the policy document.
Aditi Sharma: be.
Muddu Sudhakar: I should be able to visualize there's a policy document MD file exist. And within that policy, which policy is it violating? I should know
Aditi Sharma: So you would want to look at the policy document within this view
Muddu Sudhakar: at some point, right?

00:03:20

Muddu Sudhakar: Yeah. If I want to view visualize okay if it's violating what is the policy document I should be able to click and show that
Aditi Sharma: itself.
Muddu Sudhakar: also uh if you can if you can't do it then at least write down saying policy document and leave it there somewhere in the view I need to show that there's a policy
Aditi Sharma: Got it. Got it.
Muddu Sudhakar: document available for me to open if I want

Aditi Sharma: Okay. Okay.
Muddu Sudhakar: to
Aditi Sharma: Yeah, I can add that also. one of the feedback that I was uh like uh I was thinking of like replacing this
Muddu Sudhakar: but all these are very small fonti I don't know how you're looking at from my view everything looks
Aditi Sharma: entire okay
Muddu Sudhakar: like like I can't see it like you you have to be like some 100 by 100 vision to see it it's really small characters etc.
Aditi Sharma: I think it's because
Muddu Sudhakar: So you need to bloat up everything by at least an order couple of order of magnitude.

00:04:08

Muddu Sudhakar: Alex, can you read in text? I can't read it from where I am.
Aditi Sharma: Go.
Muddu Sudhakar: Huh?
Alex Kahng: Yeah, it's too small. It's too
Muddu Sudhakar: Yeah. I don't know why you made it so small. I mean Adi never designed it for you because you are sitting in front of it.
Alex Kahng: small.
Muddu Sudhakar: So once you projected it onto a Google Meet, I can't see anything here.
Aditi Sharma: Got it. No, I'll take that feedback. I'll change the font size. Um, what about replacing this panel with agent force? that sort of like helps you change whatever is wrong in the policy of the agent via like a conversational
Alex Kahng: So that that's okay.
Aditi Sharma: flow.
Alex Kahng: It should be you should be opening up this agent for like corrections. Basically, it should be extremely obvious that it violated something based on the policy document that should be shown. You should see you should be able to hover over a tool tip and see what policy clause was was uh breached.

00:05:08

Alex Kahng: And then basically the remediation step would be to uh you know modify in a very like scientific way uh the evals and the testing data so that it behaves better going forward. I don't think we need all the there's a lot of like strict mode block threshold but at the end of the day you just need like change the eval set or add some cases so that this clause doesn't get breached anymore. It can be very simple I think and then just close the loop there.
Aditi Sharma: Okay.
Alex Kahng: I think there should be a whole process of retraining like send it back uh for retraining with these new um you know new data points and it should retrain the model retrain the agent and then once it passes those eval it's ready for deployment. It's just that last final step in the agent uh you know development life cycle.
Aditi Sharma: Yeah, but right now you can deploy to production from here itself
Alex Kahng: Yes. But but I think like if the agent is like breaching one small thing in one small

00:06:18

Aditi Sharma: is
Alex Kahng: case toggling from strict mode to loose mode or block threshold from 90 to 95%. These are not the types of fixes you would apply to the agent.
Aditi Sharma: Okay, I probably need to do research as to what needs to be applied. I don't think
Alex Kahng: We should keep it a black box here. In my opinion,
Aditi Sharma: I
Alex Kahng: it should be not something a human will be asked to do. All the human needs to do is say add to data set, retrain agent, launch agent. That's it.
Aditi Sharma: So three call to actions add to data set retrain agent or launch agent.
Alex Kahng: Yes. And when they add to data set like it should tastefully say oh um you know here are a bunch of other examp would you like to add this data point plus a bunch of examples which we can generate based on this flavor of data point like that that sort of uh you know varied uh case addition like generation and addition to the data set should be possible here maybe with a little blowout.

00:07:33

Alex Kahng: So like you can Yeah,
Aditi Sharma: So,
Alex Kahng: go
Aditi Sharma: so when the um human says like add to data set,
Alex Kahng: ahead.
Aditi Sharma: would all the like would we see that agent is getting added added to data set here? Like would all would everything of the workflow of the agent happen here or would it happen somewhere in the
Alex Kahng: I think um like it should be a gen if you want this to be an agent interface that's actually probably
Aditi Sharma: background?
Shantanu Singh: Oops.
Alex Kahng: the best way. So if a human decides to just say add to data set, it can like create a little uh fill in a template message for the agent to for the human to send to the agent. The human can say oh add like three more data points and slightly adjust the you know the types of data points that it wants added and then the agent just streams the results. It's like okay I'm adding to the data set rerunning the evals okay making some adjustments to the agent build and uh I decided to change the threshold here now all of the evals are passing and it's in a purely like a genic manner with all the reasoning steps and and once it all the evals are passing similar to how cloud code would communicate it it would say okay everything is passing uh like deploy

00:08:47

Aditi Sharma: Yeah. Like thinking this is what I'm doing.
Alex Kahng: I think that would that as you mentioned agent force on the right side we should think cloud code on the right side but branded as agent force and yeah and the the key steps are
Aditi Sharma: Got it. Okay.
Alex Kahng: updating the evals I don't think we need to it could also update these sorts of um these sorts of uh configurations at the agent level but that should all be a black box. No human should be thinking or wasting um their creative energy on micromanaging the agent. It should be like very high level kind of.
Aditi Sharma: Got it.
Alex Kahng: Yeah.

Aditi Sharma: Okay.
Alex Kahng: Otherwise mud more or less aligned.
Muddu Sudhakar: I I think yeah I mean keep playing it. I just want to see more. Is there anything more? Oh, that's what it is.
Alex Kahng: Yeah.
Aditi Sharma: Yeah, this is what I have so far.
Muddu Sudhakar: Okay, good. At least I think we just make changes what Alex is saying.

00:09:48

Muddu Sudhakar: I think we have at least some of these things. It's getting better, guys.
Aditi Sharma: So Alex just uh I will create one workflow where like this side behaves like exactly how
Alex Kahng: Yeah.
Aditi Sharma: claude code behaves today. Okay. So basically if I tell the agent like okay add these five data sets to this agent and then run the evals then it would sort of like repolish the agent and ask you in finally that do you want to launch this agent again. Correct.
Alex Kahng: Right.
Aditi Sharma: Okay,
Alex Kahng: And it should be test driven.
Aditi Sharma: got it.
Alex Kahng: It should be eval driven. Eval first. Test first.
Shantanu Singh: Yes.
Aditi Sharma: Okay.
Alex Kahng: Yeah.
Aditi Sharma: Okay. So,
Muddu Sudhakar: So you can other thing that you can also do is evalu remember these are production agents Alex.
Aditi Sharma: um
Muddu Sudhakar: So if there are issues in the production I want to tweak it too and that's where it's going to learn also in so we want to do that feedback loop also.

00:10:33

Alex Kahng: Agreed. Here another thing we can do is we just show two sorts of interactions. One interaction is you go ahead and you see some like alert in mission control and you hone in on the agent. you retrain that agent, you come back. The second thing is maybe not an alert but uh an emergency, right? The human should be able to stop that agent or perform very drastic interactions. It's like you this is a trading system and your trading system is doing well then you can make slight slight adjustments so it's optimized but if it's about to make a trade that's very negative expected value or something a human will block it and it will be pulled in to make a manual trade. That's kind of the same sort of I guess interaction we want here where let's say an agent is in a high stakes environment which they'll eventually always be entering. you want to just pull in an expert human for the human to make the final decision. That could be what's going on here as

00:11:40

Aditi Sharma: So human in the loop but a little bit more like the story is that the agent is in a very high
Alex Kahng: well.
Aditi Sharma: stress situation
Alex Kahng: Yes. Yes.
Shantanu Singh: High state.
Alex Kahng: So yeah,
Shantanu Singh: High state. That was the story we were going towards.
Aditi Sharma: situation.
Shantanu Singh: Aditi.
Alex Kahng: that's not but it's a I think Muda's point is there's multiple interaction patterns we want here. the one about agent life cycle management and training an agent is very aligned with what Service Now is putting out with AI control tower. But when it comes to uh executing a kill switch or you know completely swapping out an agent for another agent, let's say I'm using an AWS agent and it's too expensive and I get an alert on my mission control. I'm gonna switch on over to my um my Microsoft agents.
Muddu Sudhakar: It's
Alex Kahng: I want to be able to make that switch here. Their control tower doesn't necessarily allow that. We're going to take it a step further and make this more

actionable.

00:12:37

Alex Kahng: I think Mudu was saying this is actionoriented whereas theirs is more just
Muddu Sudhakar: Yeah.
Alex Kahng: observability.
Muddu Sudhakar: So the two things I think what I understood from the conference and reading is they're still stuck on the observability side. By the couple of things have also happened if you guys saw Alex was calling Cisco just acquired another company called Galilio that's again eval.
Alex Kahng: Yeah.
Muddu Sudhakar: So I think and same thing Kronosphere is also observal.
Alex Kahng: Yeah.
Muddu Sudhakar: So Calio right? So okay take a look at that. So none of these guys are thinking through what we are thinking here. So that's good. So they'll go down that path for a while.
Alex Kahng: Good.
Muddu Sudhakar: So I think eval are good. What we want to do is in the production how do we take that input and make agents always also all these guys are missing and stopping them and security and red teaming. So basically there are three four things you should think about. One is a red teaming attack.

00:13:33

Muddu Sudhakar: Eval is just an eval then production data taking into that and evolving it. Those three aspect need to be shown and shant make
Aditi Sharma: Got it. What I'll do,
Muddu Sudhakar: sense.
Aditi Sharma: Alex, is just take the notes that you shared and I'm going to create a story for what this part of the page looks like. I just need you to approve it. Uh because the thing is lot of these terms I'm not fully aware of it. So I'm just trying to like create things from what I understand. I may lack the technical uh you know outlook for it. So if you can just give me a thumbs up, thumbs down in the Google doc itself like we did yesterday that will be very useful.

Alex Kahng: Okay, sounds
Aditi Sharma: Cool. So just finalizing the changes font needs to be changed overall.
Alex Kahng: good.
Aditi Sharma: Map needs to be more accurate and better. And then the side panel needs to be more driven of that we are going to evalu change kill all of that.

00:14:27

Alex Kahng: Right. Um, I think Yeah. Will we be able to add a second the a second set of interaction patterns? Like one like whole experience for the user would be to improve an agent and one whole experience would be to micromanage an agent and basically step in as as an agent. Are we going to offer the same interaction pattern or will we like bubble up certain actions to the user suggestions? Is it going to look the same, look and feel the same, or um
Aditi Sharma: I need to think through that.
Alex Kahng: Okay. Yeah,
Shantanu Singh: This is needed today,
Alex Kahng: because Yes,
Shantanu Singh: right? This is today is the last day for it.
Alex Kahng: that's right. I'm I'm thinking that it's sort of like um like the way people are like browser use and all these and computer use is being built out,
Shantanu Singh: So,
Alex Kahng: you can always take over control from the agent. like what does that look like in the future? Uh I think well at least for now you would basically be able to do more strategic actions such as updating and retraining your agent in the right side chat.

00:15:42

Alex Kahng: But if I'm going to micromanage the agent, I want to do so on this service graph. I want to like click in on the nodes, see its reasoning, and say micromanage like the clicks here. evaluation data set should be on the right side for for that sort of thing in in the streaming status messages that should be hidden but taking actions let's say I can see the set of actions
Aditi Sharma: Okay.
Alex Kahng: the agent was thinking about doing or the you know what were the what

was all the context surrounding the decision the bad decision that it made I want to see that in this service graph area real estate when I drill in on that agent and then I want to make a different decision decision or swap it out, right? Basically, take action over this whole system. I'm not sure exactly what it would look like, but do you understand uh what I'm saying? It's like a different interaction pattern. Or maybe you can do it via natural language on the right side, but it should be very fast and very u it like it shouldn't feel like you're sending a message into the black hole of an LLM and getting a probabilistic decision.

00:17:04

Alex Kahng: It should feel like you're really in control.
Shantanu Singh: So what is what is that you want to do with let's say this agent cluster 7 is not performing uh
Alex Kahng: Yeah, I just want to kill that agent. I want to like click the little X on the top right next to
Aditi Sharma: What if there's a little X here and you can just like close it off and just then you'll get like a model that says are you sure you want to
Alex Kahng: it.
Aditi Sharma: kill that
Shantanu Singh: So wait one by one.
Aditi Sharma: agent?
Shantanu Singh: So currently what we are doing is agent cluster 7 it's it's going rogue and what we are showing here. Can you restart and get back?
Aditi Sharma: So either you go from
Shantanu Singh: That's okay. Just go to the screen.
Aditi Sharma: Hello.
Shantanu Singh: So here we are showing this agent is not it's it has gone rogue and you know on on the right hand side you can see all the things that's not right. Now there are actions provided and we were thinking to make this right hand side panel into agent force panel and agent force is saying hey this this is the current status these are the values this is what it is not working and here are recommended actions.

00:18:07

Shantanu Singh: Now in those recommended actions, one action could be a kill
Alex Kahng: No,
Shantanu Singh: switch.
Alex Kahng: I think this should be fully conversational on the on the right side. Uh maybe.
Shantanu Singh: You see,
Alex Kahng: Yeah, may maybe. Yeah,
Shantanu Singh: so what what it it will have a followup.
Alex Kahng: it could have the kill switch maybe as a proposed action.
Shantanu Singh: I'm saying it will have a recommended action. When you do that kill switch you will have the option that how do you want to replace it. So what me and Aditi were thinking yesterday is we will ask them to replace it with a human and so agent force will give list of humans like three four humans. I will select okay for the time being this agent is avail unavailable send all the things that comes to this way to these two
Alex Kahng: Yeah. Yeah.
Shantanu Singh: humans.
Alex Kahng: I I think we should get rid of the kill switch idea. Instead, we should be able to seamlessly reroute this agent.

00:18:55

Alex Kahng: Let's say there's a graph. There's all like the edges, by the way, should map between these nodes. I should be able to reroute traffic away from this agent cluster. I should just route it visually in this left side around the agent cluster and it should say would you like to update your service graph? If I say yes and the nodes persist, all the changes persist.
Aditi Sharma: You mean like physically like pick these up and go?
Alex Kahng: Yeah.
Aditi Sharma: That is a little bit tough to do right now,
Alex Kahng: Yes. Yeah, I understand.
Aditi Sharma: Alex.
Alex Kahng: But that that's the interaction pattern I would like to see. Instead of doing it via this this graph, um instead like turn off one of these dotted lines conversationally using the right side.
Muddu Sudhakar: It's

Alex Kahng: We could do that as a workound.
Aditi Sharma: Got it.
Alex Kahng: So I can say in the right side I'll say uh stop routing any traffic to agent cluster
Aditi Sharma: Yeah.
Muddu Sudhakar: diff.
Alex Kahng: 7 or or reduce traffic to 2%.

00:19:52

Alex Kahng: Uh it seems like it's going off the rails. I'm going to make some live adjustments. Uh but during this time stop routing traffic there. And I send the message it does the thinking then it turns off the flow.
Aditi Sharma: Got
Alex Kahng: You can do something like that. Does that make sense?
Aditi Sharma: it.
Alex Kahng: Like
Shantanu Singh: Yes. So that that the idea we were building on Aditi right that we will stop this and we will send it to other
Aditi Sharma: Yeah.
Shantanu Singh: person. Now we might just say that instead of stopping completely reduce it to 2%. I feel instead of again thinking user is giving this command I I
Muddu Sudhakar: Exit.
Alex Kahng: H.
Shantanu Singh: I would go a step and say that agent force can suggest that hey this is a problem and here is what we are recommending we're recommending that stop traffic at this fix it and then for the uh you know add a human to to as an alternative for now so the first step I'll do is I'll add a human so the business doesn't stop and the traffic continues and then next step could Keep fixing

00:20:52

Aditi Sharma: I have a sorry I have a question Alex and Mudu.
Shantanu Singh: this.
Aditi Sharma: The primary use case of the user is is this an IT director who is an

operator or are they someone like what
Muddu Sudhakar: Nobody knows. Nobody knows the personas will evolve. Remember this is a futuristic.
Aditi Sharma: is
Alex Kahng: Yeah, agent admin. I think just think of it
Muddu Sudhakar: So I think how you should think of this thing like you're really curating something that you don't know how this
Alex Kahng: as
Muddu Sudhakar: will transform in the future. Imagine like these are not the existing people who's going to run this who's going to manage it. Do you think you want to give it to IT director? I don't trust that guy.
Alex Kahng: Yeah.
Muddu Sudhakar: I don't even trust anybody in IT group to be doing this like so not
Aditi Sharma: So like an
Alex Kahng: a manager.
Muddu Sudhakar: operator like we don't know it will be somebody some we have to think through which groups will run
Aditi Sharma: operator.

00:21:34

Alex Kahng: Yeah, it's going to be like a little young kid who's a savant who can put all this context of this company in his head, see it as a as an abstract view of agents and micromanage them as needed to keep this company up and running. That's it. It'll be something crazy like that.
Muddu Sudhakar: Exactly.
Alex Kahng: It's like Ender
Muddu Sudhakar: Like how look at how Elon Musk hired these young people to run this thing on
Alex Kahng: Ender.
Shantanu Singh: So Alex is describing himself I think
Muddu Sudhakar: the I
Alex Kahng: No, no, no, no,
Shantanu Singh: but this rule doesn't exist this product
Muddu Sudhakar: think
Alex Kahng: no.
Shantanu Singh: doesn't exist requirement doesn't

Aditi Sharma: I know I know I know Alex and Mudu like the thing is like because it does not exist I have zero
Shantanu Singh: exist.
Aditi Sharma: references to look at
Muddu Sudhakar: that's there's nothing I can help you with.
Aditi Sharma: right
Muddu Sudhakar: You have that's where your creativity has to come. Evangelize think That's the whole that's the whole exciting thing.

00:22:25

Aditi Sharma: Okay.
Muddu Sudhakar: You don't know what's going to happen. You you have to print in the future. You don't know who you going to use your design and the product.
Aditi Sharma: No, I got you. I got
Muddu Sudhakar: But somebody's going to somebody's buying it. Somebody's using you uh this thing already,
Aditi Sharma: you.
Muddu Sudhakar: right? So people who have to manage. Think of the future state where agents need to be managed. There will be this young people who are going to be just hired to do it and there'll be a special team
Alex Kahng: That's
Muddu Sudhakar: whose job is to not be in IT department. There'll be a separate department like you have security and IT. Will there be an agent department? I don't know what's going to
Aditi Sharma: Can you guys see this? Like one of the questions my design skill was asking that this product is a truthful observability
Muddu Sudhakar: happen
Aditi Sharma: or is it a decision theater out of which like are

00:23:08

Alex Kahng: heat.
Aditi Sharma: we
Muddu Sudhakar: observable what is in theater look it's very simple this is for let's make

practical
Aditi Sharma: you
Muddu Sudhakar: decisions today agents are all being deployed a lot of them okay many companies have deployed agents I think 90% of companies have some agent running at least one or there's utility for them. the more agents have add utility okay then people have to worry about okay now that I'm getting the value how do I manage this thing I want to observe them and I want to make sure they're doing a good job then comes cyber security for risk whenever you talk about risk and kill switch think of them like insurance policy you don't buy insurance unless you make a revenue let's say you are making zero revenue you think you're going to buy insurance policy for your life insurance for yourself no so first customers want to make money out of these agents and that money when they make money and do ROI and if it has good utility then they have to worry about hey when they when they're down I want to make sure they're doing a good job that's a happy path and happy path then I have to worry about the rogue stuff as an insurance that's evolution that will happen if you're zero revenue employee I'm not going to buy any insurance for you you're useless to me right think like that same thing for agents so right now companies are going to invest on agents to get value and utility once that utility is

00:24:32

Muddu Sudhakar: going well then they have to worry about when their downtime how do I improve their uptime like one thing you don't even talk about uptime hey which agent is down how much time what is it performance etc that's all comes under control tower that is a business aspect of the agents observability then I have to worry about rogue and misuse etc you see what I'm saying it's all evolution will happen. If there's no utility, there is no downtime, then I don't care for security. So I won't apply the security side. So think that evolution as you're telling this story. Alex,
Aditi Sharma: Got
Muddu Sudhakar: do you see what I'm saying?
Alex Kahng: Yeah.
Aditi Sharma: him.
Muddu Sudhakar: That's how I've seen this. So it's going to play out.
Alex Kahng: Yeah.
Muddu Sudhakar: How it's going to play out also? We know it. Question is who and what when the timeline we don't know.
Aditi Sharma: So just finalizing the changes. The map needs to be more accurate. The

font needs to be bigger everywhere.

00:25:29

Aditi Sharma: Whatever changes in the flow would only happen in this side panel. It will be more conversational and more agent force driven. And the story for this panel, I'm just going to create it now. In next 10 minutes, I'll share with Alex.
Muddu Sudhakar: Yeah.
Aditi Sharma: He just has to give me a thumbs up.
Muddu Sudhakar: The only other thing that he talked about is ability to make changes somewhere here. If you want to modify the agents like he talked about calling the cloud code. The thing that is missing to me is also eval are missing here. Eval should be there and also the agent production activity should be taken into
Shantanu Singh: So we are not this story does not cover evals and production.
Muddu Sudhakar: account
Shantanu Singh: We will send we will stop the agent and hand it over to human. We have the vive coding one which we are showing ev and other thing the same technological.
Muddu Sudhakar: that is during building time this is during the runtime okay yeah
Shantanu Singh: Yeah. Yeah.

00:26:12

Shantanu Singh: But it's the same technology which can be used here, right? What I'm saying
Muddu Sudhakar: but some connect the if you can take that and attach to this video is also I'm okay figure out
Shantanu Singh: is
Muddu Sudhakar: something
Alex Kahng: I think yeah I agree uh we should I mean it goes hand in hand I think here it should be abstract I mean you should just see like the the conversation part happening here you don't need to see all the internals of how it updated the data set and all that but it should stream those messages it should be a very compelling conversational experience as to you know re-eval agent on the right
Aditi Sharma: But uh Alex and Mutu given that the time that we have today right there is

limited stuff that we
Alex Kahng: side.
Muddu Sudhakar: Yeah, do that's the time comes.
Aditi Sharma: can do.
Alex Kahng: Yeah.
Muddu Sudhakar: If the time comes out, do what you the best you
Aditi Sharma: Okay.
Muddu Sudhakar: can.
Alex Kahng: One suggestion can uh as Mood said like can we blow up the font literally

00:27:03

Aditi Sharma: Sure, sure, sure, sure. I I'll do that.
Alex Kahng: uh so I like if you were to
Aditi Sharma: I'll do that. That's the first thing I'll do.
Alex Kahng: design this um can you like change your brows your screen resolution so it's like an uh for an old old uh person with no
Aditi Sharma: Yeah.
Muddu Sudhakar: Yeah,
Alex Kahng: eyesight I'm not saying
Muddu Sudhakar: he's basically saying somebody like me who is 55 years old.
Shantanu Singh: when we look at it in full screen, I think it will be better than this. But yeah, we'll increase the font overall. Just use the body size font from
Aditi Sharma: Yeah.
Alex Kahng: Because if you do that,
Aditi Sharma: Yeah.
Alex Kahng: it forces you to remove all the AI junk that gets generated here.
Shantanu Singh: Salesforce.
Aditi Sharma: Yeah. No, I'll I'll do all of that at some point. Yeah.
Alex Kahng: Like
Aditi Sharma: Um the just highlighting the story line what Shantu said that agent first chat that will help you change the agent to a human in the loop and um what was the other thing I forgot we will be able to only cover two stories.

00:28:00

Shantanu Singh: We can what we can do here is what I was thinking and then talking to Aditi is we can add a
Alex Kahng: Yeah.
Shantanu Singh: I was thinking of adding a tab that connects me to the agent builder the other demo that we had.
Alex Kahng: I see.
Muddu Sudhakar: Yeah.
Alex Kahng: I think uh Okay.
Shantanu Singh: So the way I was thinking is look at think of it like a platform where I can monitor, govern, manage agent.
Muddu Sudhakar: So
Shantanu Singh: At the same point I should be able to go and see what all agents are available and create new agents if I
Alex Kahng: Yeah, I I think that should be like that pop out icon.
Shantanu Singh: need
Alex Kahng: Anytime you hover over an agent, you can open it up in the builder or Yeah. But here in this flow, I don't think we should leave the
Muddu Sudhakar: yeah, don't put it in This chhatra don't put it here in
Shantanu Singh: Not in this not in this. How do we connect? Okay.
Muddu Sudhakar: this
Shantanu Singh: So yeah I mean that is available.

00:28:53

Shantanu Singh: So when you want to show the events and how the making the builder part of
Muddu Sudhakar: I will ask the question is can I ask you you'll give me this as a video right so I can play the video
Shantanu Singh: it
Muddu Sudhakar: tomorrow it should be a video Alex you already give me the other one as a video I want this as a
Alex Kahng: Can we
Aditi Sharma: I'll do
Alex Kahng: have you
Aditi Sharma: it.

Shantanu Singh: video is easy. Do you want video or do you want to
Alex Kahng: want the interactive site and and the video?
Aditi Sharma: I think Shant the video would be better.
Alex Kahng: Yeah.
Aditi Sharma: Do we can
Shantanu Singh: Alex just
Aditi Sharma: like Alex,
Muddu Sudhakar: Because if Alex is doing he wants a Figma one cuz Alex will be called for a
Shantanu Singh: said
Aditi Sharma: I'm going to show you a site that I made.
Muddu Sudhakar: guy.
Shantanu Singh: we will host it on Heroku. This will be a working
Alex Kahng: Yeah,
Muddu Sudhakar: Okay.
Aditi Sharma: It's already

00:29:30

Alex Kahng: that
Muddu Sudhakar: Give it to Alex.
Shantanu Singh: site.
Muddu Sudhakar: So the meeting so that you guys know 8:30 is with a person called David Ward.
Aditi Sharma: hosted.
Alex Kahng: would
Muddu Sudhakar: He's now the new chief architect for mark that he remember he made announcement he's a so I'll me and Alex will meet him for for that meeting Alex should have this whatever hero or figma my meeting with Mark I want to have video
Alex Kahng: Yeah, we need both. And worst case I can record the
Aditi Sharma: Do you want like a 11 Labs voice over in the video?
Alex Kahng: video.
Muddu Sudhakar: okay no voice or no
Aditi Sharma: For that I
Alex Kahng: You can do it.
Aditi Sharma: would.
Muddu Sudhakar: Yes.

Alex Kahng: Your voice over is way
Muddu Sudhakar: Okay.
Alex Kahng: better.
Shantanu Singh: that 11 I I find it very robotic.
Muddu Sudhakar: take your
Aditi Sharma: Okay.
Shantanu Singh: So Mud click on this click on this Aditi has update hosted it already and
Alex Kahng: Yeah. I clicked. Yeah.
Shantanu Singh: you the interactive version of the same thing what Aditi is showing and then let me know if the font you know how much increase do you want because this is how you will see it in the full

00:30:22

Aditi Sharma: Yeah.
Muddu Sudhakar: Yeah, it still should be increased. This not this small.
Shantanu Singh: browser
Alex Kahng: Let let me share my screen.
Muddu Sudhakar: Still small.
Alex Kahng: Moody, this is what you'll see uh tomorrow.
Muddu Sudhakar: Okay.
Alex Kahng: One sec.
Muddu Sudhakar: Show me your
Alex Kahng: One second.
Muddu Sudhakar: screen.
Alex Kahng: So, it'll look like uh like this right here.
Aditi Sharma: Is it like still
Alex Kahng: So, I don't know.
Aditi Sharma: small?
Alex Kahng: Is this
Shantanu Singh: I think it is. It is
Muddu Sudhakar: I think the dots of the map is confusing the heck out of me.
Alex Kahng: too
Shantanu Singh: fine.
Aditi Sharma: I will I will I will re I will redo the map.
Muddu Sudhakar: Yeah. Yeah.

Aditi Sharma: I took a
Muddu Sudhakar: I think the dots are okay.
Aditi Sharma: shortcut.
Muddu Sudhakar: That text the United States is too small at the bottom bottom right corner. Bottom right.
Aditi Sharma: Yeah. Yeah.
Muddu Sudhakar: That I can't read the Yeah.
Alex Kahng: getting cut off um because of my screen

00:31:03

Aditi Sharma: Uh actually then I'll just create map.
Alex Kahng: resolution?
Muddu Sudhakar: full.
Aditi Sharma: I did not create
Alex Kahng: So yeah, I think it's got really tight.
Aditi Sharma: one.
Muddu Sudhakar: Yeah, the LLM routing. I couldn't even read that. The inside the agent fonts are still small. Shantan know this.
Aditi Sharma: Okay, I
Muddu Sudhakar: It says corporate W AF.
Aditi Sharma: will
Muddu Sudhakar: What is this? I know WTF. What is WAF? Wall corporate. W AF is on the top left.
Shantanu Singh: Where are you reading mud?
Muddu Sudhakar: corporate a WF man here on the on top of Yeah.
Shantanu Singh: Oh, this is just a agent name or
Muddu Sudhakar: It's too small.
Aditi Sharma: Oh,
Muddu Sudhakar: I
Shantanu Singh: something.
Alex Kahng: So yeah,
Muddu Sudhakar: know.
Alex Kahng: I need to actually if you can record a video that would be helpful. I don't know the click.
Muddu Sudhakar: Live edges
Aditi Sharma: I have to move.

Muddu Sudhakar: from
Shantanu Singh: Do you have a big monitor,
Alex Kahng: Yeah.

00:31:59

Shantanu Singh: Aditi? You can put it on a big monitor and record. Everything will look bigger if it is scalable.
Aditi Sharma: Okay.
Alex Kahng: I cannot acknowledge and deploy. All right, cool.
Muddu Sudhakar: Good job though, guys. I think you're getting to We'll get some good feedback.
Alex Kahng: Yeah.
Aditi Sharma: Just um Alex,
Shantanu Singh: Yeah.
Alex Kahng: Yeah.
Aditi Sharma: next steps maybe a little bit dependency on you.
Muddu Sudhakar: It's again.
Aditi Sharma: I'm just going to send you the story. Just give me a thumbs up and thumbs down. And highlighting map change.
Alex Kahng: Yeah.
Aditi Sharma: The dots have to be aligned. Font has to be 4x. And the agent force panel would show two stories of uh human in the loop. And what was the other story? Can someone remind me? Shantu the eel.
Shantanu Singh: The evil one you you are attempting to do
Aditi Sharma: No,
Shantanu Singh: it.
Aditi Sharma: no, no. Sorry. Uh I know there was a human in the loop story.

00:32:51

Aditi Sharma: What was the second story?
Shantanu Singh: That that was it that I see there is a agent problem and agent force will give me
Aditi Sharma: Okay, cool.

Shantanu Singh: options to reduce the traffic
Muddu Sudhakar: Hey, Alex. Yeah, this guy may is waiting.
Shantanu Singh: kill.
Muddu Sudhakar: I need to go there.
Aditi Sharma: Okay. Okay.
Alex Kahng: Okay.
Aditi Sharma: Cool.
Muddu Sudhakar: Wrap up and come there because we need to have the deck to work on it.
Aditi Sharma: Thank you. Sure. Cool.
Muddu Sudhakar: Okay, I'll see you there.
Aditi Sharma: Thank you everyone.
Muddu Sudhakar: Thank you folks.
Alex Kahng: Yeah.
Muddu Sudhakar: I need to talk to Mayank on the deck as well.
Aditi Sharma: Okay. Okay.
Muddu Sudhakar: But thank you.
Shantanu Singh: Can we can we get one minute?
Muddu Sudhakar: We'll talk.
Aditi Sharma: Sure.
Shantanu Singh: Mu can we can you stay one minute? Uh Aditi Suprio,
Muddu Sudhakar: Sure.
Shantanu Singh: you can drop.
Aditi Sharma: Okay. Sure.
Alex Kahng: Okay, thanks.
Shantanu Singh: So Mudu,
Muddu Sudhakar: Tell me something.
Shantanu Singh: I wanted to if the feedback comes and we want this to

Transcription ended after 00:33:27

This editable transcript was computer generated and might contain errors. People can also change the text after it was created.